

# Bonus ConvNet MNIST training exploration

```
• begin
•     using Flux, PlutoUI, Statistics, MLDatasets, Images
•     using Flux.Data: DataLoader
•     using Flux: @epochs, onehotbatch, onecold, logitcrossentropy, throttle, unsqueeze
•     using Random: shuffle, randperm
•     using StatsBase: countmap, proportionmap
•     using IterTools:ncycle
•     using CUDA
• end
```

## Loading the Dataset

0

```
• CUDA.reclaim()
```



```
..  ..  ..  ..  ..  ..  ..  ..  ..  ..  ..  ..  ..  ..  ..  ..  ..  ..  
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0  
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0  
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0  
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0  
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0  
:   :   :   :  
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0  
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0  
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0  
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0  
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0  
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0  
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
```

```
[:, :, 10000] =  
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0  
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0  
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0  
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0  
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0  
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0  
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0  
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0  
:   :   :  
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0  
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0  
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0  
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0  
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0  
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
```

```
• begin  
•     train_x_org_, train_y_i = MNIST.traindata(Float32);  
•     test_x_org_, test_y_i = MNISTtestdata(Float32);  
• end
```



# Defining the loss and accuracy functions

loss (generic function with 1 method)

```
• function loss(x, y)
•     y_predict = model(x)
•     return Flux.Losses.crossentropy(y_predict, y)
• end
```

accuracy (generic function with 1 method)

```
• function accuracy(x, y)
•     x = Flux.onecold(model(x))
•     y = Flux.onecold(y)
•     return mean(x .== y)
• end
```

callBackEvaluation (generic function with 1 method)

```
• function callBackEvaluation()
•     lossVal = loss(test_x_org, test_y)
•     accuracyVal = accuracy(test_x_org, test_y)
•     @show(lossVal, accuracyVal)
• end
```

# Defining the Optimizer

ADADelta(0.9, IdDict())

```
• begin
•     parameters = Flux.params(model)
•     opt = Flux.Optimise.ADADelta()
• end
```

# Loading Data and Training the Model

```
DataLoader((28×28×1×10000 CuArray{Float32, 4}:
```

1

```

..  ..  ..  ..  ..  ..  ..  ..  ..  ..  ..  ..  ..  ..  ..  ..  ..  ..
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0
...   :    ..  ..  ..  ..  ..  ..  ..  ..  ..  ..  ..  ..  ..  ..  ..  ..
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0

```

```

[:, :, 1, 10000] =
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0

```

```

• begin
•     train_loader = Flux.Data.DataLoader((train_x_org, train_y), batchsize=50,
•                                         shuffle=true)
•     test_loader = Flux.Data.DataLoader((test_x_org, test_y), batchsize=50,
•                                         shuffle=true)
• end

• Flux.Optimise.train!(
•     loss,
•     parameters,
•     ncycle(train_loader, 10),
•     opt,
•     cb=Flux.throttle(callBackEvaluation, 10)
• )

```

## Model Accuracy

0.9877

- accuracy(test\_x\_org, test\_v)

## Comment:

The accuracy of our model in julia is very close to the provided tensorflow example. the difference of less than .3% in accuracy is negligible.

## Question 2

### Shuffling the each image pixel

it is worth noting that all images are being shuffled the same exact way. The pix\_perm variable is randomized but once it is compiled it is essentially a constant map of where pixels should be that when applied to an image it would do the same kind of computation.

- pix\_perm=reshape(CartesianIndices((28,28))`randperm(28\*28)], 28,28);

28×28×1×60000 CuArray{Float32, 4}:

```
[ :, :, 1, 1 ] =  
    0.0      0.0      0.0      0.992157 ... 0.0      0.0      0.0705882  0.0  
    0.0      0.0      0.0      0.192157 ... 0.0      0.0      0.0      0.0  
    0.0      0.0      0.0      0.0      ... 0.0      0.305882  0.0      0.0  
    0.0      0.0      0.992157 0.0      ... 0.0      0.0      0.419608  0.545098  
    0.0      0.0      0.0      0.0      ... 0.0      0.0      0.0      0.0  
    0.0      0.0      0.0      0.0      ... 0.0      0.0      0.0      0.0  
    0.0      0.0      0.0      0.0      ... 0.0      0.0      0.0      0.0  
    0.0      0.0      0.992157 0.0      ... 0.0      0.0      0.0      0.0  
    :           :           :           :           :           :           :           :  
    0.0      0.945098  0.0      0.0      ... 0.0      0.0      0.0      0.0  
0.141176  0.67451   0.0      0.215686 ... 0.0      0.0      0.0      0.0  
    0.0      0.0      0.0      0.0      ... 0.0      0.0      0.0      0.0  
0.509804  0.0      0.0      0.0      ... 0.0      0.713726  0.0      0.0  
    0.0      0.0      0.0      0.0      ... 0.0      0.0      0.0      0.0  
    0.0      0.0      0.0      0.0      ... 0.0      0.0      0.0      0.0  
    0.0      0.0      0.0      0.0      ... 0.0      0.0      0.0      0.992157  
  
[ :, :, 1, 2 ] =  
    0.0      0.0      0.0      0.988235 ... 0.0      0.0      0.109804  0.0  
    0.0      0.0      0.0980392 0.0      ... 0.0      0.0      0.0      0.0  
    0.0      0.0      0.0      0.0      ... 0.0      0.219608  0.0      0.0  
    0.0      0.196078  0.988235 0.0      ... 0.0      0.0      0.2      0.913725  
    0.0      0.0      0.647059  0.0      ... 0.0      0.439216  0.0      0.0  
    0.0      0.0      0.0      0.0      ... 0.729412  0.0      0.0      0.0  
    0.0      0.0      0.0      0.0      ... 0.0      0.0      0.0      0.988235  
    :           :           :           :           :           :           :  
• begin  
•     train_x_ = similar(train_x_org_);  
•     for i=1:size(train_x_,3)  
•         train_x_[:,:,i] .= train_x_org_[:, :, i][pix_perm]  
•     end
```

```

•   train_x = reshape(train_x_, 28, 28, 1, :) |> gpu
• end

```

```

28×28×1×10000 CuArray{Float32, 4}:
[:, :, 1, 1] =
 0.0      0.0      0.0      0.0      ... 0.0      0.0      0.0      0.996078
 0.0      0.0      0.0      0.72549    0.0      0.0      0.0      0.0
 0.0      0.0      0.0      0.0      0.0      0.0      0.992157  0.0
 0.0      0.0      0.776471  0.0      0.0      0.0      0.639216  0.0
 0.0      0.0      0.0      0.0      0.0      0.0      0.0      0.0
 0.0      0.0      0.0      0.0      0.0      0.0      0.0      0.0
 0.0      0.0      0.0      0.0      ... 0.0      0.0      0.0      0.0
 0.0      0.0      0.996078  0.0      0.0      0.0      0.0      0.0
⋮          ⋮          ⋮          ⋮          ⋮          ⋮          ⋮          ⋮          ⋮
0.898039  0.0      0.0      0.0      0.0      0.0      0.0      0.0
0.0      0.0      0.0      0.0      0.0      0.0      0.0      0.0
0.0      0.0      0.0      0.0      0.0      0.0      0.0      0.0
0.984314  0.0      0.0      0.0      ... 0.996078  0.776471  0.0      0.0
0.0      0.0      0.996078  0.0      0.32549    0.0      0.0      0.0
0.666667  0.0      0.156863  0.0      0.0      0.0      0.996078  0.858824

[:, :, 1, 2] =
 0.0      0.0      0.0      0.0      ... 0.223529  0.0      0.0      0.0
 0.0      0.588235  0.784314  0.301961    0.0      0.0      0.0      0.0
 0.0      0.0      0.203922  0.0      0.0      0.65098   0.0      0.0
 0.0      0.0      0.0      0.0      0.0      0.0      0.0      0.0
 0.0      0.0      0.0      0.0980392  0.0      0.305882  0.0      0.0
 0.0      0.0      0.0      0.0      ... 0.0      0.0      0.0      0.0
 0.0      0.0      0.0      0.0      0.0      0.0      0.0      0.0
⋮          ⋮          ⋮          ⋮          ⋮          ⋮          ⋮          ⋮          ⋮
• begin
•   test_x_ = similar(test_x_org_);
•   for i=1:size(test_x_,3)
•     test_x_[:, :, i] .= test_x_org_[:, :, i][pix_perm]
•   end
•   test_x = reshape(test_x_, 28, 28, 1, :) |> gpu
• end

```

## Defining a New model to be trained with the shuffled images

```

model_shuf =
Chain(Conv((3, 3), 1=>32, relu), Conv((3, 3), 32=>64, relu), MaxPool((2, 2)), Dropout(0.2
• model_shuf = Chain(
•   Conv((3, 3), 1 => 32, pad=(1, 1), relu),
•   Conv((3, 3), 32 => 64, pad=(1, 1), relu),
•   MaxPool((2, 2)),
•   Dropout(0.25),
•   flatten,
•   Dense(12544, 128, relu),
•   Dropout(0.5),
•   Dense(128, 10),
•   softmax
• ) |> gpu

```

accuracyShuf (generic function with 1 method)

```
• function accuracyShuf(x, y)
•     x = Flux.onecold(model_shuf(x))
•     y = Flux.onecold(y)
•     return mean(x .== y)
• end
```

```
lossShuf (generic function with 1 method)
```

```
• function lossShuf(x, y)
•     y_predict = model_shuf(x)
•     return Flux.Losses.crossentropy(y_predict, y)
• end
```

```
callBackEvaluationShuf (generic function with 1 method)
```

```
• function callBackEvaluationShuf()
•     lossVal = lossShuf(test_x, test_y)
•     accuracyVal = accuracyShuf(test_x, test_y)
•     @show(lossVal, accuracyVal)
• end
```

```
ADADelta(0.9, IdDict())
```

```
• begin
•     parameters_shuf = Flux.params(model_shuf)
•     opt_shuf = Flux.Optimise.ADADelta()
• end
```

DataLoader((28×28×1×10000 CuArray{Float32, 4}:

[:, :, 1, 1] =								
0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.99607
0.0	0.0	0.0	0.72549	0.0	0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.992157	0.0
0.0	0.0	0.776471	0.0	0.0	0.0	0.0	0.639216	0.0
0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
0.0	0.0	0.996078	0.0	0.0	0.0	0.0	0.0	0.0
:			..			..		
0.898039	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
0.984314	0.0	0.0	0.0	...	0.996078	0.776471	0.0	0.0
0.0	0.0	0.996078	0.0	0.32549	0.0	0.0	0.0	0.0
0.666667	0.0	0.156863	0.0	0.0	0.0	0.996078	0.85882	
[:, :, 1, 2] =								
0.0	0.0	0.0	0.0	...	0.223529	0.0	0.0	0.0
0.0	0.588235	0.784314	0.301961	0.0	0.0	0.0	0.0	0.0
0.0	0.0	0.203922	0.0	0.0	0.0	0.65098	0.0	0.0
0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0980392	0.0	0.0	0.305882	0.0	0.0
0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
:			..			..		
0.0	0.913725	0.992157	0.0	0.0	0.0	0.0	0.0	0.0
0.992157	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0	...	0.992157	0.992157	0.0	0.0
0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.992
[:, :, 1, 3] =								
0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0
0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0
0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0
0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0
0.0	0.0	0.0	0.0	0.803922	0.0	0.0	0.0	0
0.0	0.0	0.0	0.592157	0.0	...	0.0	0.0	0
0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0
:			..			..		
0.0	0.223529	0.0	0.0	0.0	0.0	0.0	0.0	0
0.0	0.0	0.0	0.0	0.321569	0.0	0.0	0.0	0
0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0
0.0	0.0	0.0	0.0	...	0.0	0.792157	0.329412	0
0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0
0.0	0.0	0.0	0.0	0.529412	0.0	0.0	0.0	0
...								
[:, :, 1, 9998] =								
0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0
0.0	0.0	0.14902	0.0	0.0	0.0	0.0	0.0	0
0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.690196	0
0.0	0.0	0.290196	0.0	0.0	0.0	0.0	0.847059	0
0.0	0.0	0.996078	0.054902	0.0	0.901961	0.0	0.0	0
0.0	0.0	0.0	0.690196	...	0.0	0.0	0.0	0
0.0	0.0	0.996078	0.0	0.0	0.0	0.0	0.0	0
:			..			..		
0.996078	0.682353	0.996078	0.0	0.0	0.0	0.0	0.0	0
0.0	0.0	0.0	0.0	0.0	0.0	0.482353	0.0	0

```

    ..  ..  ..  ..  ..  ..  ..  ..
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.701961
0.0  0.996078 0.0392157 0.0  0.0  0.4  0.0  0.0
0.0  0.0  0.0  0.0  0.996078 0.0  0.0  0.0
0.0  0.0  0.0  0.0  0.0  ... 0.0  0.0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
:
0.298039 0.972549 0.996078 0.0  0.996078 0.0  0.0  0.0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
0.0  0.0  0.0  0.0  0.996078 0.0  0.0  0.0
0.996078 0.32549 0.0  0.0  0.0  ... 0.623529 0.0  0.54902
0.0  0.0  0.0  0.0  0.0  0.866667 0.0  0.0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0

```

`[:, :, 1, 10000] =`

```

0.0  0.458824 0.0  0.0  0.0  ... 0.0  0.0
0.0  0.0196078 0.811765 0.0  0.0  0.0  0.0  0.0
0.0  0.0  0.0  0.0  0.0  0.992157 0.0
0.0  0.0  0.992157 0.0  0.0  0.0  0.294118
0.0  0.0  0.980392 0.992157 0.992157 0.94902 0.0
0.0  0.0  0.0  0.992157 0.0  ... 0.0  0.0
0.0  0.992157 0.992157 0.0  0.0  0.0  0.0  0.0
:
0.0  0.0  0.25098 0.0  0.0  0.0  0.0  0.0
0.172549 0.0  0.831373 0.0  0.776471 0.992157 0.0196078
0.0  0.0  0.0  0.0  0.996078 0.0  0.0
0.337255 0.992157 0.0  0.490196 0.0  ... 0.960784 0.0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0

```

```

• begin
•   train_loader_shuf = Flux.Data.DataLoader((train_x, train_y), batchsize=50,
•     shuffle=true)
•   test_loader_shuf = Flux.Data.DataLoader((test_x, test_y), batchsize=50,
•     shuffle=true)
• end

• Flux.Optimise.train!(
•   lossShuf,
•   parameters_shuf,
•   ncycle(train_loader_shuf, 10),
•   opt_shuf,
•   cb=Flux.throttle(callBackEvaluationShuf, 10)
• )

```

## Model Accuracy

0.9719

- `accuracyShuf(test_x, test_y)`

## Observation and Comments

In this experiment, initialization is very important. When `pix_perm` compiles it sets a random transformation that will be applied to the entire dataset. And it stands to reason that sometimes this transformation is going to be more forgiving than otherwise.

In all of our experiments we have had accuracies in the high nineties. The model trained on the shuffled images was only marginally less accurate than the model trained on the original images.

We know that CNN architecture are sensitive to variability and diversity of pixel intensity within and between local regions. It is a spatially aware model. In this case we did shuffle our images but we shuffled them all the same way, still giving the model a stable pattern to learn for each class.

If we shuffled each individual picture randomly (compiled `pix_perm` in a for loop before every shuffle for example) the accuracy would have suffered greatly.

## Question 3

### Defining a model with no dropout

```
ADADelta(0.9, IdDict())  
• begin  
•     model_no_dropout = Chain(  
•         Conv((3, 3), 1 => 32, pad=(1, 1), relu),  
•         Conv((3, 3), 32 => 64, pad=(1, 1), relu),  
•         MaxPool((2, 2)),  
•         flatten,  
•         Dense(12544, 128, relu),  
•         Dense(128, 10),  
•         softmax  
•     ) |> gpu  
•  
•     function accuracy_no_dropout(x, y)  
•         x = Flux.onecold(model_no_dropout(x))  
•         y = Flux.onecold(y)  
•         return mean(x .== y)  
•     end  
•  
•     function loss_no_dropout(x, y)  
•         y_predict = model_no_dropout(x)  
•         return Flux.Losses.crossentropy(y_predict, y)  
•     end
```

```
•     parameters_no_dropout = Flux.params(model_no_dropout)
•     opt_no_dropout = Flux.Optimise.ADADelta()
• end
```

## This is the modek that we made in question 1 (optimal dropout)

```
ADADelta(0.9, IdDict())
begin
    model_dropout = Chain(
        Conv((3, 3), 1 => 32, pad=(1, 1), relu),
        Conv((3, 3), 32 => 64, pad=(1, 1), relu),
        MaxPool((2, 2)),
        Dropout(0.25),
        flatten,
        Dense(12544, 128, relu),
        Dropout(0.5),
        Dense(128, 10),
        softmax
    ) |> gpu
    function accuracy_dropout(x, y)
        x = Flux.onecold(model_dropout(x))
        y = Flux.onecold(y)
        return mean(x .== y)
    end
    function loss_dropout(x, y)
        y_predict = model_dropout(x)
        return Flux.Losses.crossentropy(y_predict, y)
    end
    parameters_dropout = Flux.params(model_dropout)
    opt_dropout = Flux.Optimise.ADADelta()
end
```

## Defining a model with high dropout

```
ADADelta(0.9, IdDict())
begin
    model_high_dropout = Chain(
        Conv((3, 3), 1 => 32, pad=(1, 1), relu),
        Conv((3, 3), 32 => 64, pad=(1, 1), relu),
        MaxPool((2, 2)),
        Dropout(0.6),
        flatten,
        Dense(12544, 128, relu),
        Dropout(0.8),
        Dense(128, 10),
        softmax
```

```
•     ) |> gpu
•
•     function accuracy_high_dropout(x, y)
•         x = Flux.onecold(model_dropout(x))
•         y = Flux.onecold(y)
•         return mean(x .== y)
•     end
•
•     function loss_high_dropout(x, y)
•         y_predict = model_dropout(x)
•         return Flux.Losses.crossentropy(y_predict, y)
•     end
•
•     parameters_high_dropout = Flux.params(model_dropout)
•     opt_high_dropout = Flux.Optimise.ADADelta()
• end
```

customTrainingLoop (generic function with 1 method)

```
•     function customTrainingLoop(model_cust, trainLoader, test_x, test_y)
•         function loss_cust(x, y)
•             y_predict = model_cust(x)
•             return Flux.Losses.crossentropy(y_predict, y)
•         end
•
•         function accuracy_cust(x, y)
•             x = Flux.onecold(model_cust(x))
•             y = Flux.onecold(y)
•             return mean(x .== y)
•         end
•
•         opt_cust = Flux.Optimise.ADAM()
•         params_cust = Flux.params(model_cust)
•         losses = []
•         accuracies = []
•         for (train_point, test_point) in trainLoader
•             grads = gradient(params_cust) do
•                 loss_cust(train_point, test_point)
•             end
•             Flux.Optimise.update!(opt_cust, params_cust, grads)
•             append!(losses, loss_cust(test_x, test_y))
•             append!(accuracies, accuracy_cust(test_x, test_y))
•             @show accuracies
•             if accuracies[end] > 0.95
•                 break
•             end
•         end
•         return losses, accuracies
•     end
• 
```

## We will train the three models

Each model will be trained for a single epoch and record the loss and the accuracy for each model

```
([2.26851, 2.20679, 2.13904, 2.06445, 1.97705, 1.87709, 1.75624, 1.61533, 1.48362,
```

```
• dropout_losses, dropout_accuracies = customTrainingLoop(model_dropout, train_loader,
```

### test\_x\_org. test\_v)

```
([2.28677, 2.27866, 2.25622, 2.23074, 2.2038, 2.17881, 2.16029, 2.14272, 2.12386,
```

- `high_dropout_losses, high_dropout_accuracies = customTrainingLoop(model_high_dropout, train_loader. test_x_org. test_v)`

```
([2.16151, 2.01287, 1.83101, 1.65656, 1.50071, 1.36136, 1.13703, 0.945228, 0.93651
```

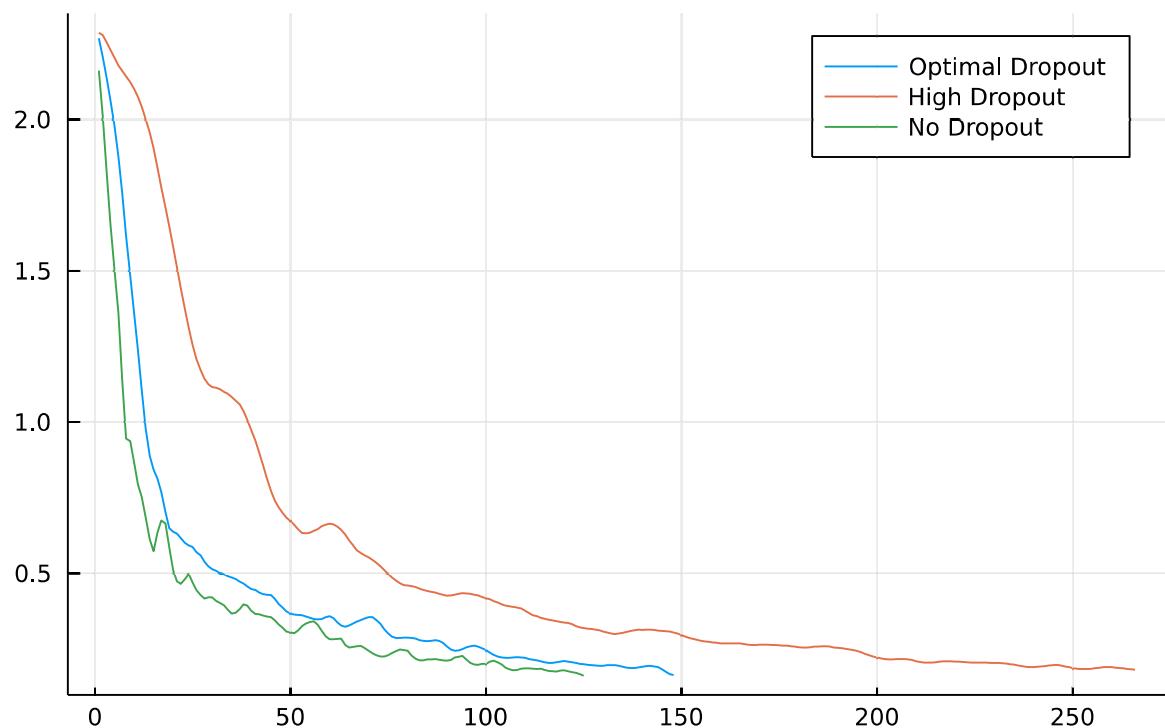
- `no_dropout_losses, no_dropout_accuracies = customTrainingLoop(model_no_dropout, train_loader. test_x_org. test_v)`

- `using Plots`

`list_comp` (generic function with 1 method)

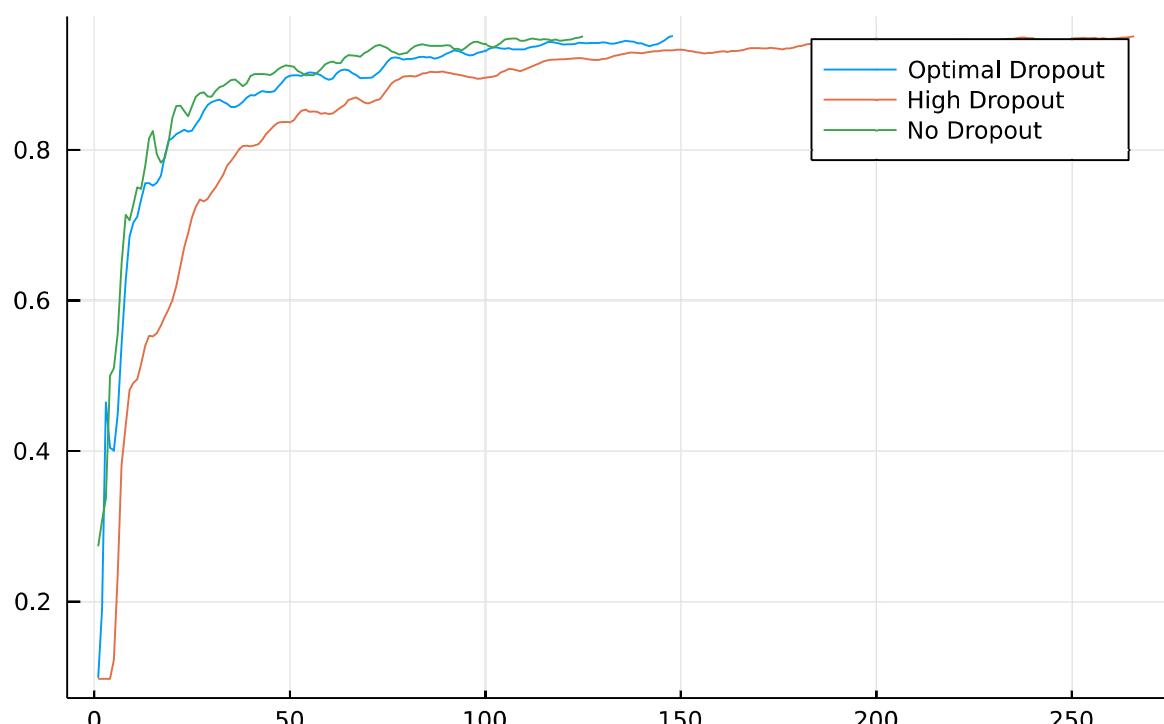
- `function list_comp(list_)`
- `return [x for x in 1:length(list_)]`
- `end`

Losses



- `begin`
- `plot(list_comp(dropout_losses), dropout_losses, label="Optimal Dropout", title="Losses")`
- `plot!(list_comp(high_dropout_losses), high_dropout_losses, label="High Dropout")`
- `plot!(list_comp(no_dropout_losses), no_dropout_losses, label="No Dropout")`
- `end`

## Accuracies



```
• begin
•     plot(list_comp(dropout_accuracies), dropout_accuracies, label="Optimal Dropout",
      title="Accuracies")
•     plot!(list_comp(high_dropout_accuracies), high_dropout_accuracies, label="High
      Dropout")
•     plot!(list_comp(no_dropout_accuracies), no_dropout_accuracies, label="No
      Dropout")
• end
```

## Comments

To truly measure whether our model is over fitting, we would have to withhold a portion of the dataset for validation, some data that the model never sees during training to test the fully trained mode. An Overfit model will perform poorly on the validation data despite having a very high training accuracy. We couldn't achieve this in this experiment due to constraints related to my machine's gpu memory and the fact that Pluto seems to run the whole sheet and overload the memory everytime a change is made.

Therefore we will try to estimate the model's tendency to overfit by how fast the model training accuracy improves. we will use early stoping to compare the losses and accuracies of the three models.

We notice that the model with no dropout at all was the fastest to reach 95% accuracy after just about 120 batches (of 50) of the pictures. This model is likely to overfit during the rest of the training period because of how often it will get exposed to similar data. The model might suffer from over-

reliance on a few of its inputs.

We can trust the dropout models to generalize better because they get rid of this over-reliance. Since not all of the models inputs are present at all time during training, the model does not become over-reliant on its inputs and as a result can generalize better.

## Assignment Questions

### Question 1 (10+ pts)

Replicate the Conv net described in [here](#) at In [20] in Julia. Do you get similar results?  
Comment!

### Question 2 (5+ pts)

After shuffling the pixels around, did the results change? Comment.

### Question 2 (5+ pts)

Would the Conv net *overfits* faster or slower without the dropout layers? Demonstrate with results and comment. Get even more points with experimentation and analysis of different dropout rates.