

# Fashion MNIST training

```
PlotlyBackend()
```

```
1 begin
2     using Markdown
3     using InteractiveUtils
4     using Flux, PlutoUI, Statistics, MLDatasets, Images
5     using Flux.Data: DataLoader
6     using Flux: @epochs, onehotbatch, onecold, logitcrossentropy, throttle, unsqueeze
7     using Random: shuffle
8     using StatsBase: countmap, proportionmap
9     using IterTools: ncycle
10    using CUDA
11    using Plots
12    Plots.PlotlyBackend()
13 end
```

```
1 begin
2     using DarkMode
3     DarkMode.enable(theme="monokai")
4 end
```

```
784×10000 CuArray{Float32, 2}:
0.0 0.0      0.0 0.0 0.0      0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0      0.0 0.0 0.0      0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0      0.0 0.0 0.0      0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0      0.0 0.0 0.00784314 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0      0.0 0.0 0.00392157 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0      0.0 0.0 0.00392157 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0      0.0 0.0 0.00392157 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
⋮           ⋮   ⋮           ⋮
0.0 0.682353 0.0 0.0 0.0      0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.741176 0.0 0.0 0.0      0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.262745 0.0 0.0 0.0      0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0      0.0 0.0 0.0      0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0      0.0 0.0 0.0      0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0      0.0 0.0 0.0      0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
```

We can look at few samples. Note that I do `1. .- tensor` to change the background to white

I use the above table to construct a dictionary to make labels to descriptions. Stored in the `class_dict` variable

```
classDict =
Dict(0 => "T-shirt/top", 4 => "Coat", 5 => "Sandal", 6 => "Shirt", 2 => "Pullover",
```

# Question 1: Defining the model

```
model = Chain(Dense(784, 300, relu), Dense(300, 100, relu), Dense(100, 10), softmax)

1 model = Chain(
2     Flux.Dense(784, 300, relu),
3     Flux.Dense(300, 100, relu),
4     Flux.Dense(100, 10),
5     softmax
6 ) ⚡ gpu
```

# Question 2

## Defining Loss Function

```
loss (generic function with 1 method)

1 function loss(x, y)
2     y_predict = model(x)
3     return Flux.Losses.crossentropy(y_predict, y)
4 end
```

## Defining Accuracy Function

```
accuracy (generic function with 1 method)

1 function accuracy(x, y)
2     x = Flux.onecold(model(x))
3     y = Flux.onecold(y)
4     return mean(x .== y)
5 end
6
```

## Instantiate an ADAM optimizer

```
opt = ⚡ADAM(0.001, (0.9, 0.999), IdDict())

1 opt = Flux.Optimise.ADAM()
```

# Question 3

# Defining badIdx Function

We will start by defining a function called getClass; this function is supposed to return the class name of a test datapoint or a test datalabel. We ended up having to define two methods for the same function to handle both test datapoints and labels.

```
getClass (generic function with 1 method)
```

```
1 function getClass(tensorIn, model)
2     classIx = Flux.onecold(model(tensorIn), 0:9)
3     return classDict[classIx]
4 end
```

```
getClass (generic function with 2 methods)
```

```
1 function getClass(tensorIn::Flux.OneHotArray{UInt32,10,0,1,UInt32})
2     classIx = Flux.onecold(tensorIn, 0:9)
3     return classDict[classIx]
4 end
```

We used Our implementation of getClass to construct our badIdx function

```
badIdx (generic function with 1 method)
```

```
• function badIdx(test_x, test_y, model)
•     n = size(test_x)[2]
•     doesNotMatch = []
•     for i in 1:n
•         trueClass = getClass(test_y[:,i])
•         if getClass(test_x[:,i], model) != trueClass
•             push!(doesNotMatch, i)
•         end
•     end
•     return doesNotMatch
• end
```

# Question 4

Creating the model parameters

```
parameters =
Params([Float32[-0.06551085 0.0715978 ... -0.04103929 0.018368045; -0.03336854 0.008918841
• parameters = Flux.params(model)
```

Creating dataloaders for the test and training datasets

```
■ DataLoader((784×10000 CuArray{Float32, 2}):
  0.0  0.0      0.0  0.0  0.0      0.0 ...  0.0  0.0  0.0  0.0  0.0  0.0
  0.0  0.0      0.0  0.0  0.0      0.0 ...  0.0  0.0  0.0  0.0  0.0  0.0
  0.0  0.0      0.0  0.0  0.0      0.0 ...  0.0  0.0  0.0  0.0  0.0  0.0
  0.0  0.0      0.0  0.0  0.00784314  0.0 ...  0.0  0.0  0.0  0.0  0.0  0.0
  0.0  0.0      0.0  0.0  0.0      0.0 ...  0.0  0.0  0.0  0.0  0.0  0.0
  0.0  0.0      0.0  0.0  0.00392157  0.0 ...  0.0  0.0  0.0  0.0  0.0  0.0
  0.0  0.0      0.0  0.0  0.00392157  0.0 ...  0.0  0.0  0.0  0.0  0.0  0.0
  ...
  :       :       :       :       :
  0.0  0.682353  0.0  0.0  0.0      0.0 ...  0.0  0.0  0.0  0.0  0.0  0.0
  0.0  0.741176  0.0  0.0  0.0      0.0 ...  0.0  0.0  0.0  0.0  0.0  0.0
  0.0  0.262745  0.0  0.0  0.0      0.0 ...  0.0  0.0  0.0  0.0  0.0  0.0
  0.0  0.0      0.0  0.0  0.0      0.0 ...  0.0  0.0  0.0  0.0  0.0  0.0
  0.0  0.0      0.0  0.0  0.0      0.0 ...  0.0  0.0  0.0  0.0  0.0  0.0
  0.0  0.0      0.0  0.0  0.0      0.0 ...  0.0  0.0  0.0  0.0  0.0  0.0
```

```
• begin
•   train_loader = Flux.Data.DataLoader((train_x, train_y), batchsize=50,
  shuffle=true)
•   test_loader = Flux.Data.DataLoader((test_x, test_y), batchsize=50, shuffle=true)
• end
```

Defining a callback function, which prints to cunsole the loss and accuracy on the validation dataset when triggered

```
• md"Defining a callback function, which prints to cunsole the loss and accuracy on the validation dataset when triggered"
```

callBackEvaluation (generic function with 1 method)

```
• function callBackEvaluation()
•   lossVal = loss(test_x, test_y)
•   accuracyVal = accuracy(test_x, test_y)
•   @show(lossVal, accuracyVal)
• end
```

Training the Model and Using ncycle and DataLoader. We specified that the callback funtion is to be triggered every 10 seconds

```
• Flux.Optimise.train!(
•   loss,
•   parameters,
•   ncycle(train_loader, 10),
•   opt,
•   cb=Flux.throttle(callBackEvaluation, 10)
• )
```

## Question 5: Calculating Classes Error Rates

We defined a function called getClassMismatchCount; which uses a given model to classify the validation dataset, and then compares its predictions to the given validation labels.

The function will output the number of misclassified pictures for each class, sorted by from worst to best performing

```
getClassMismatchCount (generic function with 1 method)
• function getClassMismatchCount(test_x, test_y, model)
•     n = size(test_y)[2]
•     classMismatchCount = Dict()
•     predictions = Flux.onecold(model(test_x))
•     for i in 1:n
•         trueClass = getClass(test_y[:,i])
•         if classDict[predictions[i]-1] != trueClass
•             if haskey(classMismatchCount, trueClass)
•                 classMismatchCount[trueClass] += 1
•             else
•                 classMismatchCount[trueClass] = 1
•             end
•         end
•     end
•     mismatchCountVector = collect(classMismatchCount)
•     return sort(mismatchCountVector, by=x -> x[2], rev=true)
• end
```

Next we created a function called getClassErrorRates, which uses the output from getClassMismatchCount and devides each class misclassified picture count by the number of datapoints in the class. the function returns the error rate for each class sorted from worst to best performing

```
classCounts =
Dict("T-shirt/top" => 1000, "Trouser" => 1000, "Shirt" => 1000, "Coat" => 1000, "Sandals" => 1000, "Pants" => 1000, "Sandal/Slipper" => 1000, "Dress" => 1000, "Trousers" => 1000, "Skirt" => 1000)
• classCounts = Dict([(classDict[i].count(x -> x == i) / test_v_i) for i in 0:9])
```

```
getClassErrorRates (generic function with 1 method)
• function getClassErrorRates(mismatchCountVector)
•     classErrorRates = Dict()
•     for pair in mismatchCountVector
•         classErrorRates[pair[1]] = pair[2] / classCounts[pair[1]]
•     end
•     classErrorRatesVector = collect(classErrorRates)
•     return sort(classErrorRatesVector, by=x -> x[2], rev=true)
• end
```

Finally we defined a function called analyzeMisclassification, this function will tell us for each class and identifies what wrong classes the model had predicted for it, and for each of these wrong classes, the function will calculate the percentage of the respective wrong class compared to the total number of misclassification.

In this way we will understand how the model misclassifies some items of clothing as others.

```
analyzeMisclassification (generic function with 1 method)
• function analyzeMisclassification(test_x, test_y, model)
•     n = size(test_y)[2]
•     classMismatch = Dict()
```

```
•     classMismatchCount = Dict()
•     predictions = Flux.onecold(model(test_x))
•     for i in 1:n
•         trueClass = getClass(test_y[:,i])
•         predictedClass = classDict[predictions[i]-1]
•         if predictedClass != trueClass
•             if ~haskey(classMismatch, trueClass)
•                 classMismatch[trueClass] = Dict()
•                 classMismatchCount[trueClass] = 1
•             else
•                 classMismatchCount[trueClass] += 1
•             end
•             if haskey(classMismatch[trueClass], predictedClass)
•                 classMismatch[trueClass][predictedClass] += 1
•             else
•                 classMismatch[trueClass][predictedClass] = 1
•             end
•         end
•     end
•     for (key, _) in classMismatch
•         for (key2, value) in classMismatchCount
•             if haskey(classMismatch[key], key2)
•                 classMismatch[key][key2] /= classMismatchCount[key]
•             end
•         end
•         classMismatch[key] = collect(classMismatch[key])
•         classMismatch[key] = sort(classMismatch[key], by=x -> x[2], rev=true)
•     end
•     return classMismatch
• end
```

Testing getClass

```
prediction = "Ankle boot"
• prediction = getClass(test_x[1]. model)

actual = "Ankle boot"
• actual = getClass(test_v[1])
```

Calculating the final accuracy of the MLP network

```
finalAccuracy = 0.1252
• finalAccuracy = accuracy(test_x, test_v)
```

running badIdx on the results of the MLP

```
badIndices = □Any[  
    1: 13  
    2: 18  
    3: 24  
    4: 26  
    5: 41  
    6: 43  
    7: 44  
    8: 50  
    9: 52  
    10: 58  
    11: 67  
    12: 68  
    13: 69  
    14: 90  
    15: 99  
    16: 128  
    17: 146  
    18: 148  
    19: 151  
    20: 152  
  
    more  
  
    1143: 9942  
    1144: 9947  
    1145: 9948  
    1146: 9954  
    1147: 9956  
    1148: 9962  
    1149: 9973  
    1150: 9978  
    1151: 9980  
    1152: 9992  
]
```

```
• badIndices = badIdx(test_x, test_v, model)
```

```
classMisMatchCounts = □Pair{Any, Any}[  
    1: "Shirt" ⇒ 353  
    2: "Coat" ⇒ 282  
    3: "Pullover" ⇒ 152  
    4: "T-shirt/top" ⇒ 108  
    5: "Dress" ⇒ 71  
    6: "Sneaker" ⇒ 70  
    7: "Trouser" ⇒ 39  
    8: "Sandal" ⇒ 30  
    9: "Ankle boot" ⇒ 30  
    10: "Bag" ⇒ 17  
]
```

```
• classMisMatchCounts = getClassMismatchCount(test_x, test_v, model)
```

Running getClassMismatchCount on the results of the MLP.

We see that the top 3 misclassified classes are in order:

1. Shirt at 35.3% accuracy
2. Coat at 28.2% accuracy
3. Pullover at 15.2% accuracy

```
classMisMatchRates = Pair{Any, Any}[
    1: "Shirt" => 0.353
    2: "Coat" => 0.282
    3: "Pullover" => 0.152
    4: "T-shirt/top" => 0.108
    5: "Dress" => 0.071
    6: "Sneaker" => 0.07
    7: "Trouser" => 0.039
    8: "Sandal" => 0.03
    9: "Ankle boot" => 0.03
    10: "Bag" => 0.017
]
• classMisMatchRates = getClassErrorRates(classMisMatchCounts)
```

## Question 6: Analyzing Class Error Rates

---

We ran analyzeMisclassification on the results of the MLP to understand the mistakes that the model made

```
classMisMatchAnalysis =  
Dict{Any, Any}(  
    "Sneaker" => Pair{Any, Any}[  
        1: "Ankle boot" => 0.757143  
        2: "Sandal" => 0.242857  
    ]  
    "Shirt" => Pair{Any, Any}[  
        1: "T-shirt/top" => 0.447592  
        2: "Pullover" => 0.303116  
        3: "Coat" => 0.133144  
        4: "Dress" => 0.0736544  
        5: "Bag" => 0.0424929  
    ]  
    "Trouser" => Pair{Any, Any}[  
        1: "Dress" => 0.589744  
        2: "T-shirt/top" => 0.205128  
        3: "Pullover" => 0.0769231  
        4: "Coat" => 0.0512821  
        5: "Bag" => 0.0512821  
        6: "Shirt" => 0.025641  
    ]  
    "Coat" => Pair{Any, Any}[  
        1: "Pullover" => 0.510638  
        2: "Shirt" => 0.234043  
        3: "Dress" => 0.198582  
        4: "Bag" => 0.0531915  
        5: "T-shirt/top" => 0.0035461  
    ]  
    "T-shirt/top" => Pair{Any, Any}[  
        1: "Shirt" => 0.555556  
        2: "Dress" => 0.25  
        3: "Pullover" => 0.0925926  
        4: "Bag" => 0.0740741  
        5: "Sandal" => 0.0185185  
        6: "Coat" => 0.00925926  
    ]  
    "Sandal" => Pair{Any, Any}[  
        1: "Sneaker" => 0.533333  
        2: "Ankle boot" => 0.466667  
    ]  
    "Bag" => Pair{Any, Any}[  
        1: "T-shirt/top" => 0.294118  
        2: "Sandal" => 0.235294  
        3: "Dress" => 0.235294  
        4: "Sneaker" => 0.117647  
        5: "Shirt" => 0.0588235  
        6: "Coat" => 0.0588235  
    ]  
    "Dress" => Pair{Any, Any}[  
        1: "T-shirt/top" => 0.338028  
        2: "Shirt" => 0.211268  
        3: "Coat" => 0.183099  
        4: "Pullover" => 0.140845  
        5: "Bag" => 0.0985915  
        6: "Trouser" => 0.028169  
    ]  
    "Pullover" => Pair{Any, Any}[  
        1: "Coat" => 0.447368  
        2: "Shirt" => 0.276316  
        3: "T-shirt/top" => 0.171053  
        4: "Dress" => 0.0723684  
        5: "Bag" => 0.0328947  
    ]
```

```
• classMisMatchAnalvsis = analvzeMisclassification(test_x, test_v, model)
```

Here we defined some functions to do plots that we need to visualize the model results

PlotPair (generic function with 1 method)

```
• begin
•     function PlotPair(pairs, label, title)
•         names = []
•         values= []
•         for pair in pairs
•             push!(names, pair[1])
•             push!(values, pair[2])
•         end
•         p = bar(values, xticks=(1:10, names), label=label, title=title, xrotation=45)
•         return p
•     end
• end
```

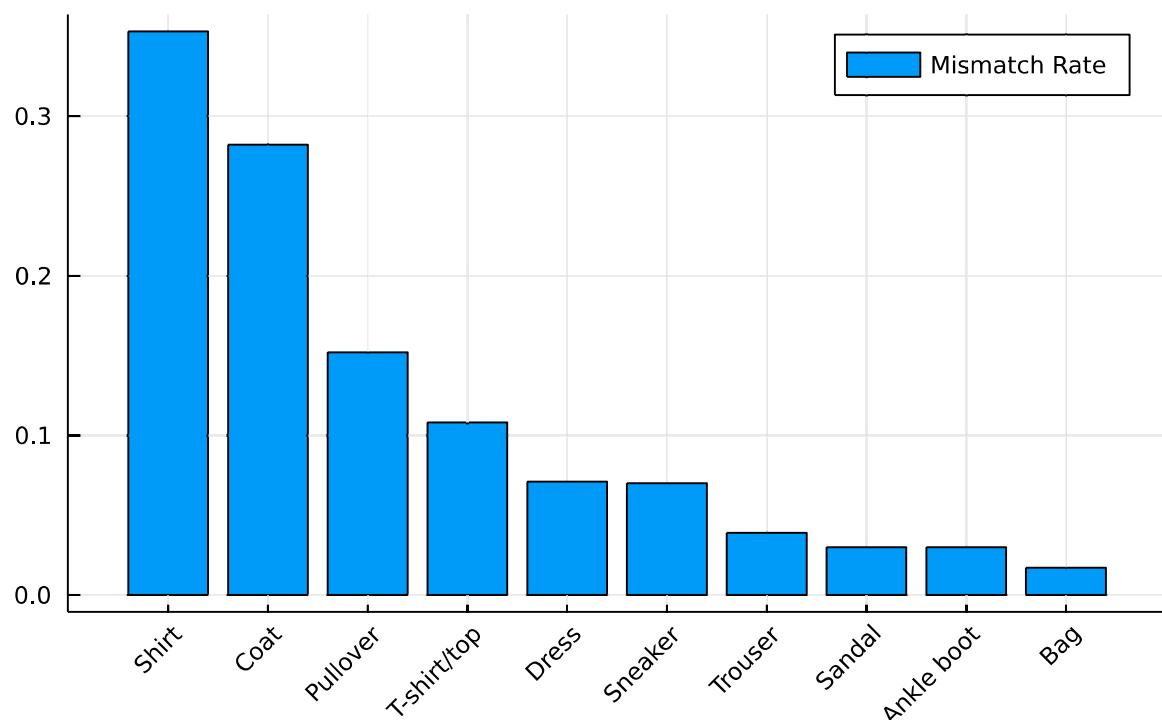
PlotDict (generic function with 1 method)

```
• begin
•     function PlotDict(classMisMatchAnalysis)
•         nFigures = length(classMisMatchAnalysis)
•         l = @layout [grid(ceil(nFigures/2) |> Int, 2)]
•         bars = []
•         for (key, value) in classMisMatchAnalysis
•             names = []
•             values= []
•             for pair in value
•                 push!(names, pair[1])
•                 push!(values, pair[2])
•             end
•             bar_ = bar(values, xticks=(1:10, names), label="mismatch rate",
•             title=key, xrotation=45)
•             push!(bars, bar_)
•         end
•         p = plot(bars..., layout = (ceil(nFigures/2) |> Int, 2), size=(1000,1500))
•         return p
•     end
• end
```

in this plot we visualized the error rate for each class. We notice that Shirt, Coat and T-shirt/top are the top three misclassified classes.

We also notice that articles of clothing that are worn on the upper body e.g. Shirt, Coat, T-shirt/top, pullover and dress, have much higher error rates compared to other classes, suggesting that they are often misclassified as each other. We will analyze that further shortly.

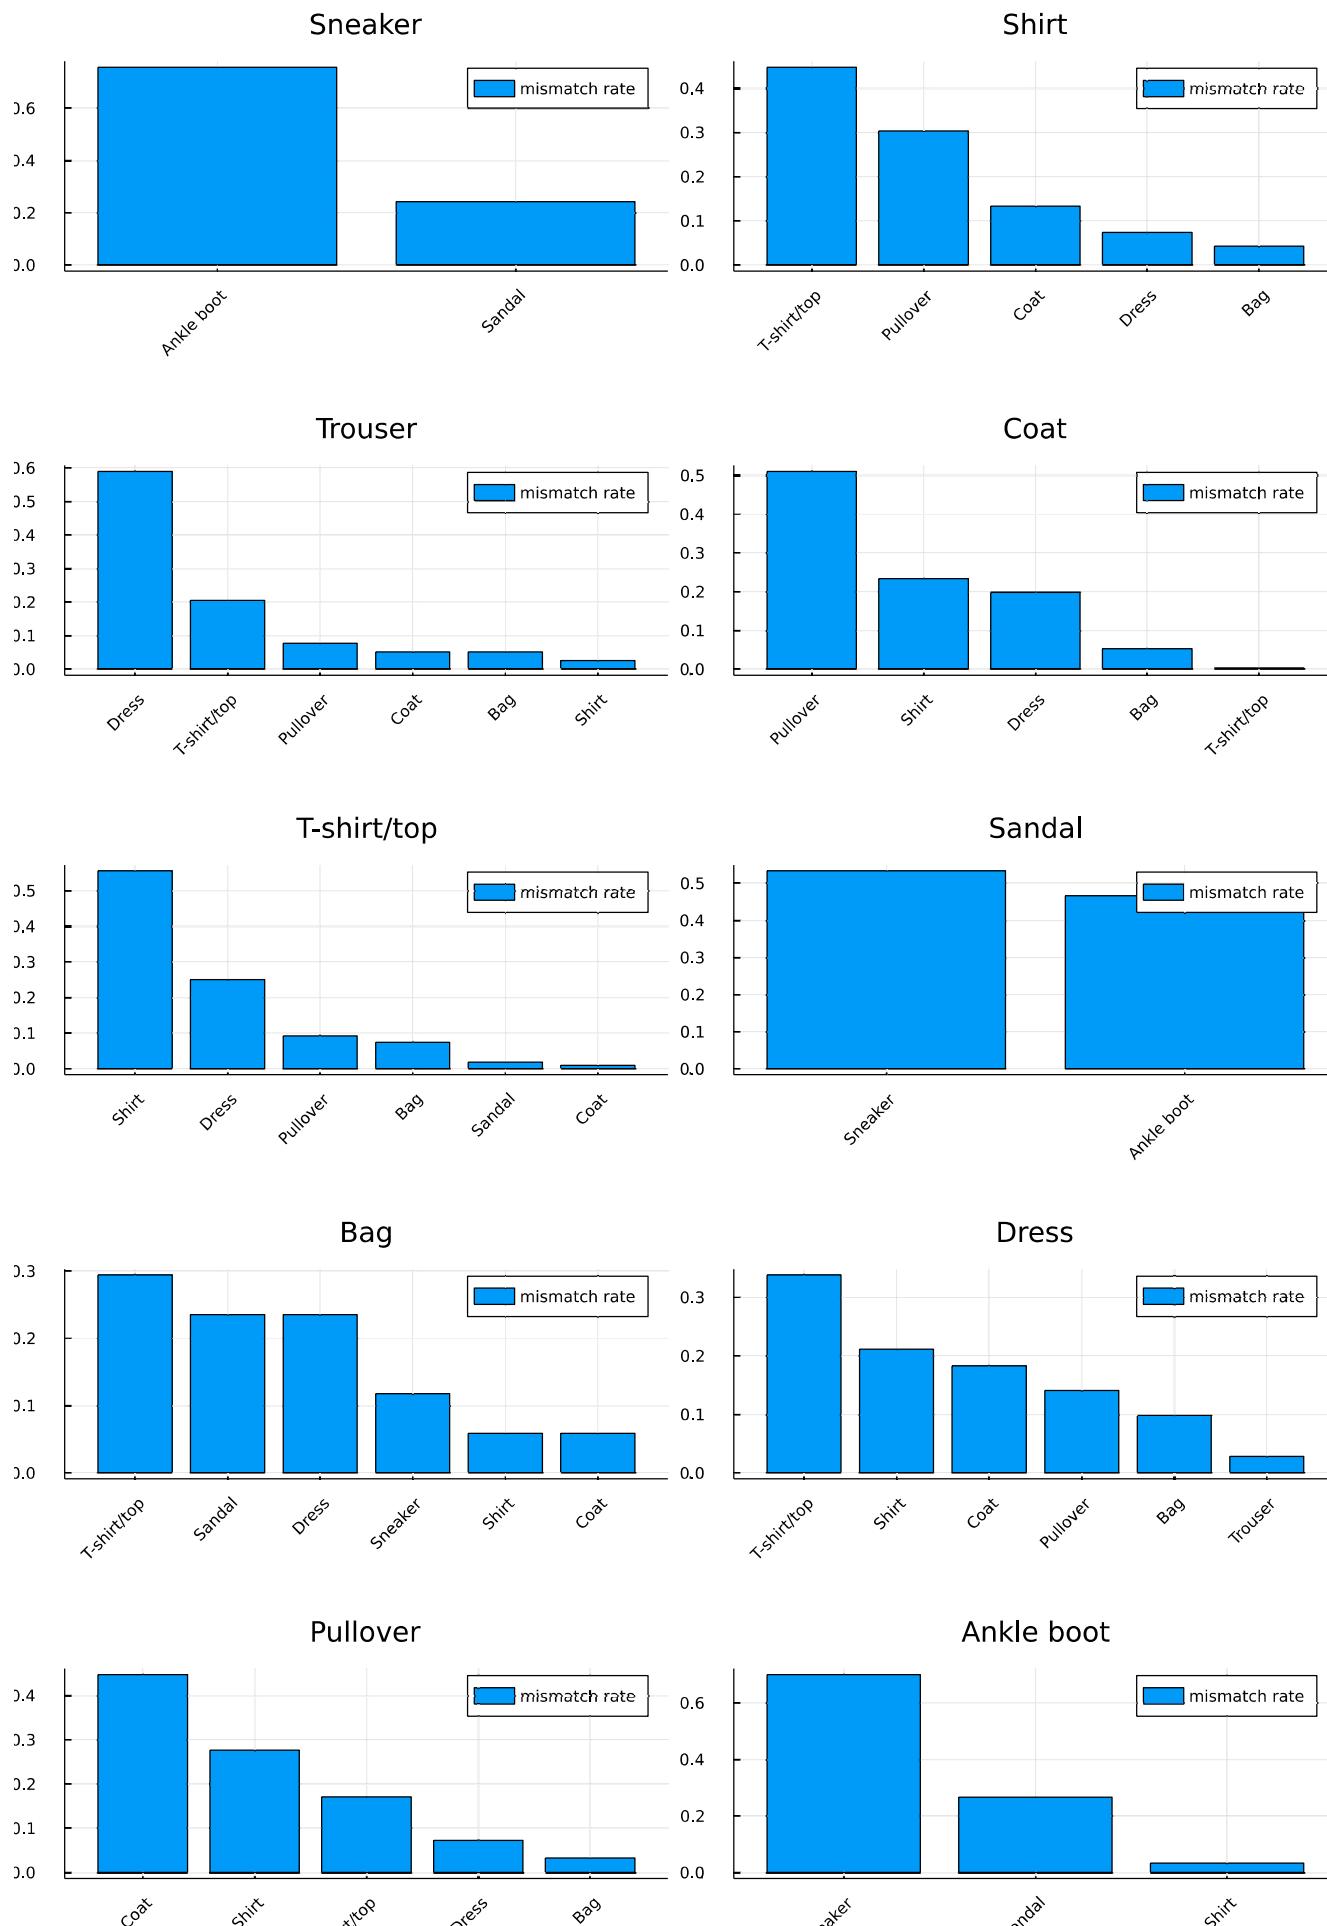
## Error Rate Per Class



```
• PlotPair(classMisMatchRates, "Mismatch Rate", "Error Rate Per Class")
```

In this plot we are looking at each individual class in the validation dataset, and identifying what percentage of its misclassification instances were mapped to what wrong class.

For example, We see that the misclassified sneakers, were identified as either sandals or ankle boots by the model.. etc.



- `PlotDict(classMisMatchAnalvsis)`

# Question 6 (repeated): Building a Custom Training Loop

- `md"># Question 6 (repeated): Building a Custom Training Loop"`

We made an identical MLP model to train with the custom training loop

```
model_cust =  
Chain(Dense(784, 300, relu), Dense(300, 100, relu), Dense(100, 10), softmax)  
  
• model_cust = Chain(  
•     Flux.Dense(784, 300, relu), # Flattened input Image -> 300 Nodes Hidden Layer 1  
•     Flux.Dense(300, 100, relu), # 300 Nodes Hidden Layer 1-> 100 Nodes Hidden Layer 2  
•     Flux.Dense(100, 10), # 100 Nodes Hidden Layer 2-> 10 Nodes Output  
•     softmax # Softmax for Classification  
• ) |> gpu  
  
accuracy_cust (generic function with 1 method)  
• function accuracy_cust(x, y, model_cust)  
•     x = Flux.onecold(model_cust(x))  
•     y = Flux.onecold(y)  
•     return mean(x .== y)  
• end
```

Defining the custom training loop. This loop will train a model for 10 epochs and will output the validation loss and the accuracy after each epoch.

```
customTrainingLoop (generic function with 1 method)
```

```
• function customTrainingLoop(model_cust, trainLoader, test_x, test_y)
•     function loss_cust(x, y)
•         y_predict = model_cust(x)
•         return Flux.Losses.crossentropy(y_predict, y)
•     end
•
•     opt_cust = Flux.Optimise.ADAM()
•     params_cust = Flux.params(model_cust)
•
•     for _ in 1:10
•         for (train_point, test_point) in trainLoader
•             grads = gradient(params_cust) do
•                 loss_cust(train_point, test_point)
•             end
•             Flux.Optimise.update!(opt_cust, params_cust, grads)
•         end
•         @show loss_cust(test_x, test_y), accuracy_cust(test_x, test_y, model_cust)
•     end
• end
```

training the second MLP with the custom training loop and getting its final accuracy

```
• customTrainingLoop(model_cust, train_loader, test_x, test_v)
```

0.8839

```
• accuracy_cust(test_x, test_v, model_cust)
```

## Question 7: Creating a CNN Model

For the CNN model we have to reshape the dataset before we can start training the model

```
train_x_conv =  
28×28×1×60000 CuArray{Float32, 4}:  
[:, :, 1, 1] =  
 0.0 0.0 0.0 0.0      0.0      ... 0.0      0.00784314 0.0 0.0 0.0  
 0.0 0.0 0.0 0.0      0.0      ... 0.0      0.0 0.0 0.0  
 0.0 0.0 0.0 0.0      0.0      ... 0.290196 0.0 0.0 0.0 0.0  
 0.0 0.0 0.0 0.0      0.0      ... 0.741176 0.0 0.0 0.0 0.0  
 0.0 0.0 0.0 0.0      0.0      ... 0.831373 0.258824 0.0 0.0 0.0  
 0.0 0.0 0.0 0.0      0.0      ... 0.74902 0.784314 0.0 0.0 0.0  
 0.0 0.0 0.0 0.0      0.0      ... 0.686275 0.870588 0.0 0.0 0.0  
 ...  
 0.0 0.0 0.0 0.0      0.00392157 0.737255 0.65098 0.0 0.0 0.0  
 0.0 0.0 0.0 0.0      0.0117647 0.760784 0.658824 0.0 0.0 0.0  
 0.0 0.0 0.0 0.0      0.0156863 0.752941 0.388235 0.0 0.0 0.0  
 0.0 0.0 0.0 0.00392157 0.0 0.847059 0.227451 0.0 0.0 0.0  
 0.0 0.0 0.0 0.00392157 0.0 0.666667 0.0 0.0 0.0 0.0  
 0.0 0.0 0.0 0.0      0.0117647 0.0 0.0 0.0 0.0 0.0  
[:, :, 1, 2] =  
 0.0      0.0      0.0      ... 0.0      0.0 0.0 0.0  
 0.0      0.0      0.0      ... 0.0      0.0 0.0 0.0  
 0.0      0.0      0.0      ... 0.0      0.0 0.0 0.0  
 0.0      0.00392157 0.0      0.0      0.0 0.0 0.0 0.0  
 0.0      0.0      0.0      ... 0.0      0.0 0.0 0.0  
 0.00392157 0.0      0.054902 ... 0.00392157 0.0 0.0 0.00392157  
 0.0      0.0      0.690196 0.0      0.0 0.0 0.0 0.0  
 • train_x_conv = reshape(train_x_. 28. 28. 1. :) |> gpu
```

```
train_loader_conv =
```

```
■DataLoader((28×28×1×60000 CuArray{Float32, 4}:
```

```
[:, :, 1, 1] =
 0.0 0.0 0.0 0.0      0.0      ... 0.0      0.00784314 0.0 0.0
 0.0 0.0 0.0 0.0      0.0      ... 0.0      0.0      0.0 0.0
 0.0 0.0 0.0 0.0      0.0      ... 0.290196 0.0      0.0 0.0
 0.0 0.0 0.0 0.0      0.0      ... 0.741176 0.0      0.0 0.0
 0.0 0.0 0.0 0.0      0.0      ... 0.831373 0.258824 0.0 0.0
 0.0 0.0 0.0 0.0      0.0      ... 0.74902 0.784314 0.0 0.0
 0.0 0.0 0.0 0.0      0.0      ... 0.686275 0.870588 0.0 0.0
 ⋮
 0.0 0.0 0.0 0.0      0.00392157 0.737255 0.65098 0.0 0.0
 0.0 0.0 0.0 0.0      0.0117647 0.760784 0.658824 0.0 0.0
 0.0 0.0 0.0 0.0      0.0156863 0.752941 0.388235 0.0 0.0
 0.0 0.0 0.0 0.00392157 0.0      ... 0.847059 0.227451 0.0 0.0
 0.0 0.0 0.0 0.00392157 0.0      ... 0.666667 0.0      0.0 0.0
 0.0 0.0 0.0 0.0      0.0117647 0.0      0.0      0.0 0.0
```

```
[:, :, 1, 2] =
```

```
0.0      0.0      0.0      ... 0.0      0.0 0.0 0.0
 0.0      0.0      0.0      ... 0.0      0.0 0.0 0.0
 0.0      0.0      0.0      ... 0.0      0.0 0.0 0.0
 0.0      0.00392157 0.0      ... 0.0      0.0 0.0 0.0
 0.0      0.0      0.0      ... 0.0      0.0 0.0 0.0
 0.00392157 0.0      0.054902 ... 0.00392157 0.0 0.0 0.003921
 0.0      0.0      0.690196 0.0      ... 0.0      0.0 0.0 0.0
 ⋮
 0.0      0.176471 0.878431 0.0      0.0 0.0 0.0 0.0
 0.0      0.0      0.643137 0.00392157 0.0 0.00784314 0.0
 0.0      0.0      0.0      ... 0.0      0.0 0.0 0.0
 0.0      0.0      0.0      ... 0.0      0.0 0.0 0.0
 0.0      0.0      0.0      ... 0.0      0.0 0.0 0.0
 0.0      0.0      0.0      ... 0.0      0.0 0.0 0.0
```

```
[:, :, 1, 3] =
```

```
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0
 ⋮
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0
```

```
...
```

```
[:, :, 1, 59998] =
```

```
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0
 ⋮
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0
```

- `train_loader_conv = Flux.Data.DataLoader((train_x_conv, train_y), batchsize=50, shuffle=true)`

```

test_x_conv =
28×28×1×10000 CuArray{Float32, 4}:
[:, :, 1, 1] =
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0470588 0.0 0.0 0.0 0.0 0.0 0.0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.262745 0.0 0.0 0.0 0.0 0.0 0.0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.415686 0.0 0.0 0.0 0.0 0.0 0.0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.643137 0.0 0.0 0.0 0.0 0.0 0.0
 ⋮ ⋮ ⋮ ⋮ ⋮ ⋮ ⋮ ⋮ ⋮ ⋮ ⋮ ⋮ ⋮ ⋮ ⋮ ⋮ ⋮ ⋮
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.737255 0.0 0.0 0.0 0.0 0.0 0.0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.686275 0.0 0.0 0.0 0.0 0.0 0.0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.635294 0.0 0.0 0.0 0.0 0.0 0.0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.619608 0.0 0.0 0.0 0.0 0.0 0.0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.592157 0.0 0.0 0.0 0.0 0.0 0.0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0431373 0.0 0.0 0.0 0.0 0.0 0.0

```

```

[:, :, 1, 2] =
 0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.737255 0.509804 0.298039 0.192157
 0.0 0.0 0.0 0.4 0.913725 ... 0.941176 0.960784 1.0 0.803922
 0.0 0.0 0.537255 1.0 0.945098 0.980392 0.94902 0.976471 0.772549

```

- `test_x_conv = reshape(test_x_. 28. 28. 1. :) |> gpu`

Here we defined a CNN model as described in the assignment questions

```

model_conv =
Chain(Conv((5, 5), 1=>6, relu), MaxPool((2, 2)), Conv((5, 5), 6=>16, relu), MaxPool((2, 2))

```

- `model_conv = Chain(`
- `Conv((5, 5), 1=>6, pad=SamePad(), relu),`
- `MaxPool((2, 2)),`
- `Conv((5, 5), 6=>16, pad=SamePad(), relu),`
- `MaxPool((2, 2)),`
- `flatten,`
- `Dense(784, 10),`
- `softmax,`
- `) |> gpu`

We trained the CNN model with our custom training loop. And tested its final accuracy

- `customTrainingLoop(model_conv. train_loader_conv. test_x_conv. test_v)`

0.9

- `accuracy_cust(test_x_conv. test_v. model_conv)`

## Question 8: Improving the CNN Model Through Architectural Changes

```

model_conv_2 =
Chain(Conv((3, 3), 1=>16, relu), MaxPool((2, 2)), Conv((3, 3), 16=>32, relu), MaxPool((2,
• model_conv_2 = Chain(
•     # First convolution, operating upon a 28x28 image
•     Conv((3, 3), 1 => 16, pad=(1, 1), relu),
•     MaxPool((2, 2)),
•
•     # Second convolution, operating upon a 14x14 image
•     Conv((3, 3), 16 => 32, pad=(1, 1), relu),
•     MaxPool((2, 2)),
•
•     # Third convolution, operating upon a 7x7 image
•     Conv((3, 3), 32 => 32, pad=(1, 1), relu),
•     MaxPool((2, 2)),
•
•     # 'Flux.flatten' will make the pictures (3, 3, 32, N)
•     flatten,
•     Dense(288, 128, relu),
•     Dense(128, 64, relu),
•     Dense(64, 10),
•     softmax
• ) |> gpu

```

We train the new CNN model with our custom training loop

```

• customTrainingLoop(model_conv_2, train_loader_conv, test_x_conv, test_v)

```

We can see that this model performs better at more than 90% accuracy

0.9122

```

• accuracy_cust(test_x_conv, test_v, model_conv_2)

```

We will do the same analysis we did on the results of the first MLP model

```

classMisMatchCountsConv = □Pair{Any, Any}[
    1: "Shirt" => 250
    2: "T-shirt/top" => 157
    3: "Coat" => 141
    4: "Pullover" => 126
    5: "Dress" => 83
    6: "Ankle boot" => 45
    7: "Sneaker" => 23
    8: "Trouser" => 18
    9: "Sandal" => 18
    10: "Bag" => 17
]

```

```

• classMisMatchCountsConv = getClassMismatchCount(test_x_conv, test_v, model_conv_2)

```

We again see that the top misclassified classes are all upperbody clothing items. however the top three misclassified items here are different:

1. Shirt at 25% error Rate.

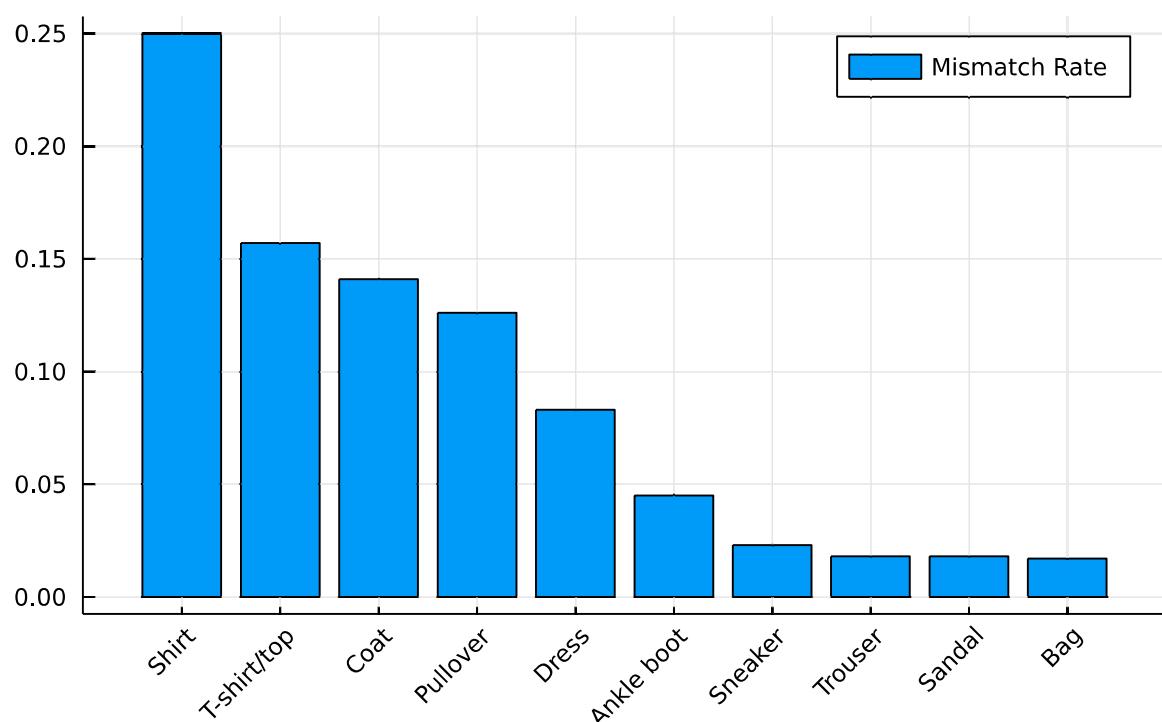
2. T-shirt/top at 15.7% error Rate.

3. Coat 14.1% error Rate.

```
classMisMatchRatesConv = □Pair{Any, Any}[
    1: "Shirt" ⇒ 0.25
    2: "T-shirt/top" ⇒ 0.157
    3: "Coat" ⇒ 0.141
    4: "Pullover" ⇒ 0.126
    5: "Dress" ⇒ 0.083
    6: "Ankle boot" ⇒ 0.045
    7: "Sneaker" ⇒ 0.023
    8: "Trouser" ⇒ 0.018
    9: "Sandal" ⇒ 0.018
    10: "Bag" ⇒ 0.017
]
```

```
• classMisMatchRatesConv = getClassErrorRates(classMisMatchCountsConv)
```

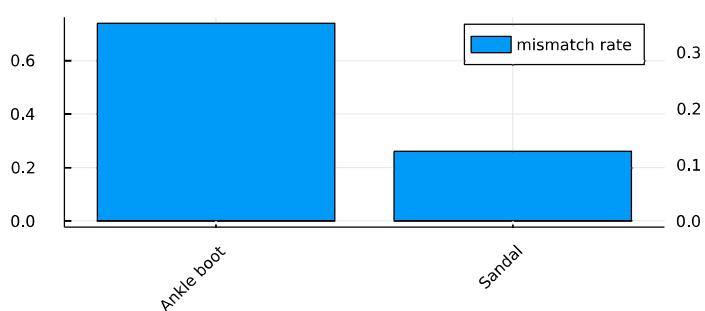
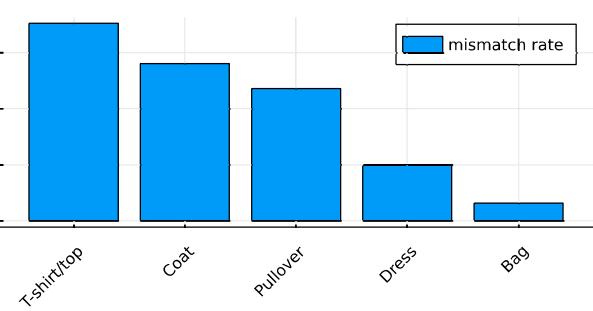
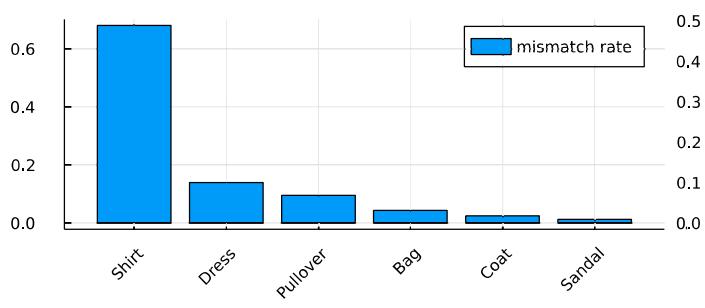
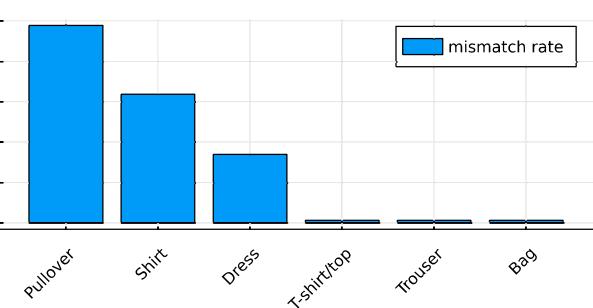
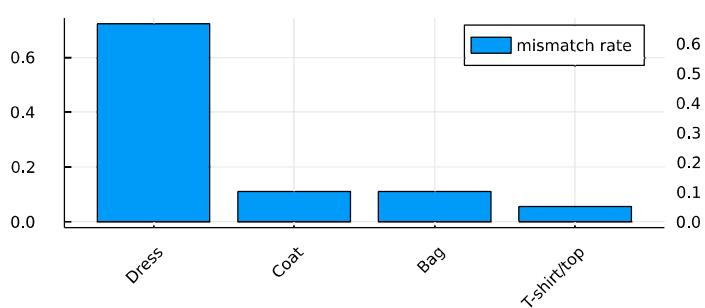
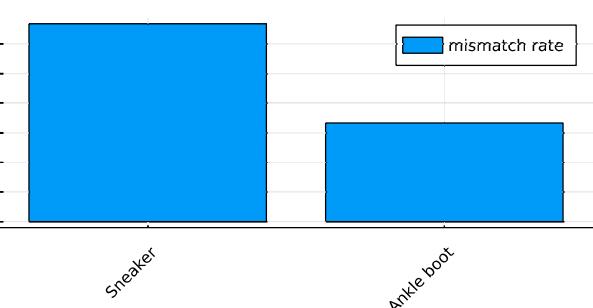
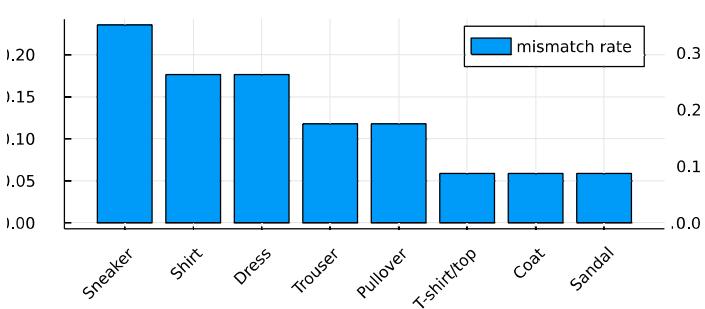
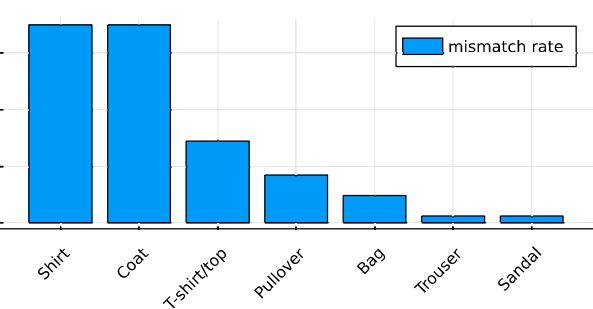
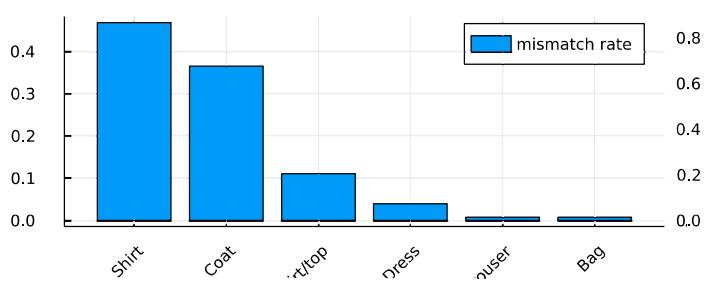
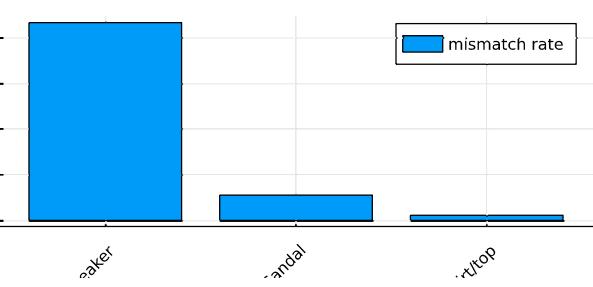
Error Per Class



```
• PlotPair(classMisMatchRatesConv, "Mismatch Rate", "Error Per Class")
```

```
classErrorAnalysisConv2 =  
Dict{Any, Any}(  
    "Sneaker" => Pair{Any, Any}[  
        1: "Ankle boot" => 0.73913  
        2: "Sandal" => 0.26087  
    ]  
    "Shirt" => Pair{Any, Any}[  
        1: "T-shirt/top" => 0.352  
        2: "Coat" => 0.28  
        3: "Pullover" => 0.236  
        4: "Dress" => 0.1  
        5: "Bag" => 0.032  
    ]  
    "T-shirt/top" => Pair{Any, Any}[  
        1: "Shirt" => 0.681529  
        2: "Dress" => 0.140127  
        3: "Pullover" => 0.0955414  
        4: "Bag" => 0.044586  
        5: "Coat" => 0.0254777  
        6: "Sandal" => 0.0127389  
    ]  
    "Coat" => Pair{Any, Any}[  
        1: "Pullover" => 0.489362  
        2: "Shirt" => 0.319149  
        3: "Dress" => 0.170213  
        4: "T-shirt/top" => 0.0070922  
        5: "Trouser" => 0.0070922  
        6: "Bag" => 0.0070922  
    ]  
    "Trouser" => Pair{Any, Any}[  
        1: "Dress" => 0.722222  
        2: "Coat" => 0.111111  
        3: "Bag" => 0.111111  
        4: "T-shirt/top" => 0.0555556  
    ]  
    "Sandal" => Pair{Any, Any}[  
        1: "Sneaker" => 0.666667  
        2: "Ankle boot" => 0.333333  
    ]  
    "Bag" => Pair{Any, Any}[  
        1: "Sneaker" => 0.235294  
        2: "Shirt" => 0.176471  
        3: "Dress" => 0.176471  
        4: "Trouser" => 0.117647  
        5: "Pullover" => 0.117647  
        6: "T-shirt/top" => 0.0588235  
        7: "Coat" => 0.0588235  
        8: "Sandal" => 0.0588235  
    ]  
    "Dress" => Pair{Any, Any}[  
        1: "Shirt" => 0.349398  
        2: "Coat" => 0.349398  
        3: "T-shirt/top" => 0.144578  
        4: "Pullover" => 0.0843373  
        5: "Bag" => 0.0481928  
        6: "Trouser" => 0.0120482  
        7: "Sandal" => 0.0120482  
    ]  
    "Pullover" => Pair{Any, Any}[  
        1: "Shirt" => 0.468254  
        2: "Coat" => 0.365079  
        3: "T-shirt/top" => 0.111111  
    ]
```

```
• classErrorAnalvsisConv2 = analvzeMisclassification(test_x_conv, test_v, model_conv_2)
```

**Sneaker****Shirt****T-shirt/top****Coat****Trouser****Sandal****Bag****Dress****Pullover****Ankle boot**

- `PlotDict(classErrorAnalvsisConv2)`

# Question 9

yes, the top 3 mmisclassified claasses changed across the models we trained, they also changed if we trained the same model

## Assignment Questions

### Question 1 (4 pts)

Using Flux Chain build a Dense MLP NN, with following layers in sequence 784, 300, 100, 10. Run the output of the lastad layer through a softmax function.

### Question 2 (2 pts)

- Define a loss funtion based on crossentropy
- Instantiate an ADAM optimizer
- Define an accuracy function that could runn over the whole dataset (Hint: make use of onecold and mean)

### Question 3 (3 pts)

Define a funciton badIdx that filter for the indices where the model failes to classify correctly.

### Question 4 (3 pts)

Make use of nycle and DataLoader and Flux.train! to train the model with batchsize of 50 and 10 epochs (Hint: see ). Use a callback function, cb to the accuracy and loss every 10 seconds.

### Question 5 (4 pts)

Which top 3 classes did the Dense MLP NN model struggle with the most? Are ther error rates uniform across the classes? If the repeat the training, it the top three misclassified classes

3×3×20×1

### Question 6 (4 pts)

Which top 3 classes did the Dense MLP NN model struggle with the most? Are the error rates uniform across the classes? If you repeat the training, will the top three misclassified classes

### Question 6 (3 pts)

Build your own *custom training loop* with same parameters. Check that it works just like in question 4.

### Question 7 (5 pts)

Construct a convolutional model with following architecture

1. Conv with a (5,5) kernel mapping to \*\*6\*\* feature maps, with 'relu', same padding
2. 2x2 Max pool
3. Conv with a (5,5) kernel mapping to \*\*16\*\* feature maps, with 'relu', same padding
4. 2x2 Max pool
5. Appropriately sized Dense layer with 10 outputs
6. A 'softmax' layer

Do the training using a *custom training loop*

### Question 8 (3 pts)

Do changes in the convnet architecture to get a better accuracy.

### Question 9 (3 pts)

Did the top 3 classes in Question 6 change with the use of the networks in Questions 8 or 9?